

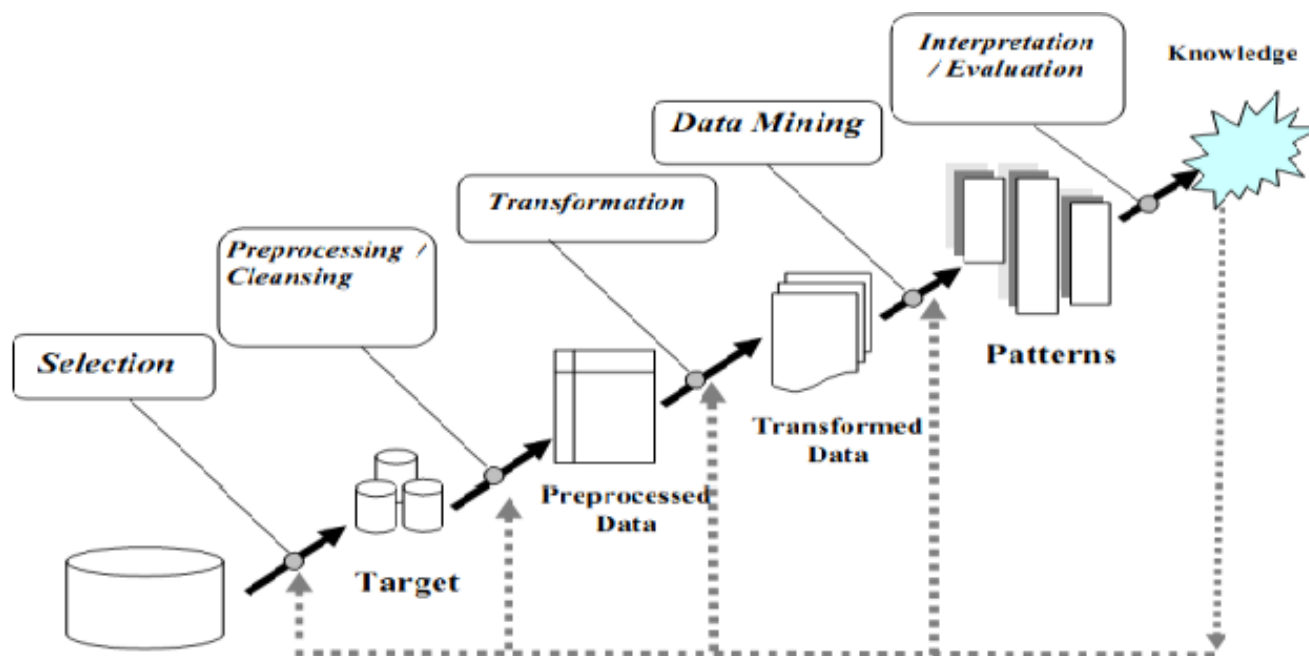
Data Analytics Process Approaches

Data Analytics Approaches

- Generic
 - KDD
 - CRISP-DM
- Vendor specific
 - Google
 - Microsoft
 - Cloudera


Knowledge Discovery for Databases (KDD)

- First implemented in 1996
- Divides the process for finding knowledge in data



CRISP-DM

- Published shortly after KDD
- Active EU project in 1997 with an effort to update the process in 2006–2008
- Major contributor was Daimler-Benz
- Described as an iterative approach and methodology to solve data-mining problems

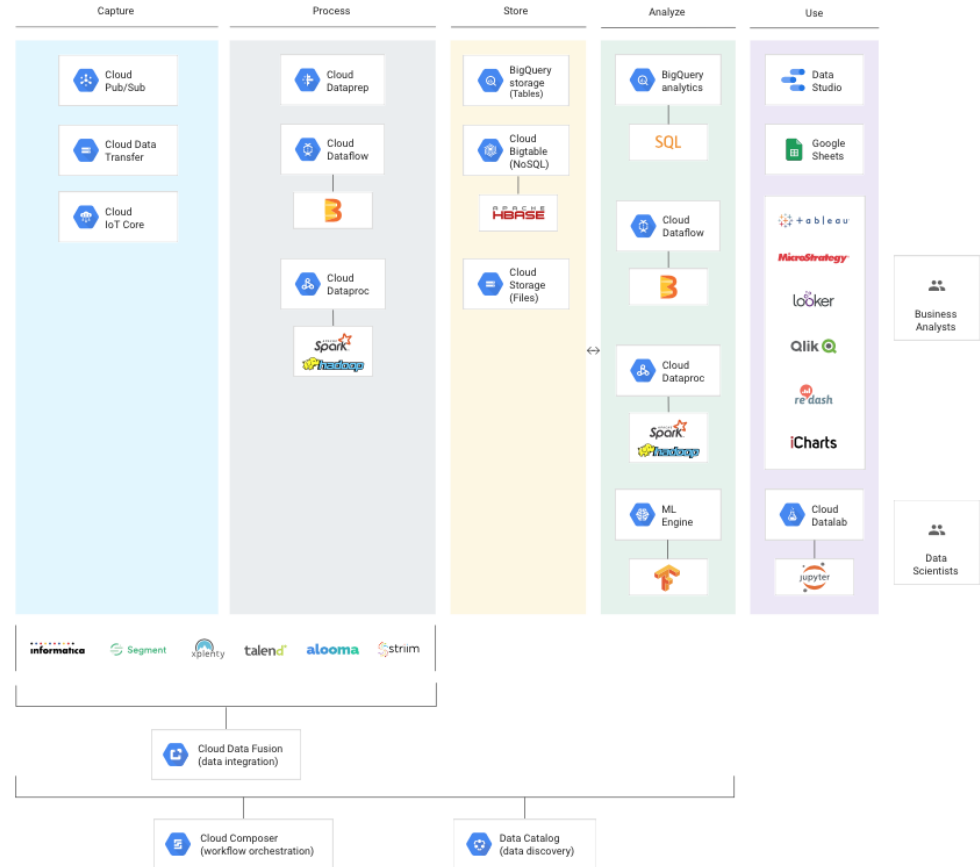


Vendor Approaches

- Cloud ML players
 - Google
 - Amazon Web Services (AWS)
 - Microsoft
 - Cloudera




























Google Approaches

- Google Cloud's big data solutions page (<https://cloud.google.com/solutions/big-data>)
- Steps
 - Capture
 - Process
 - Store
 - Analyze
 - Use



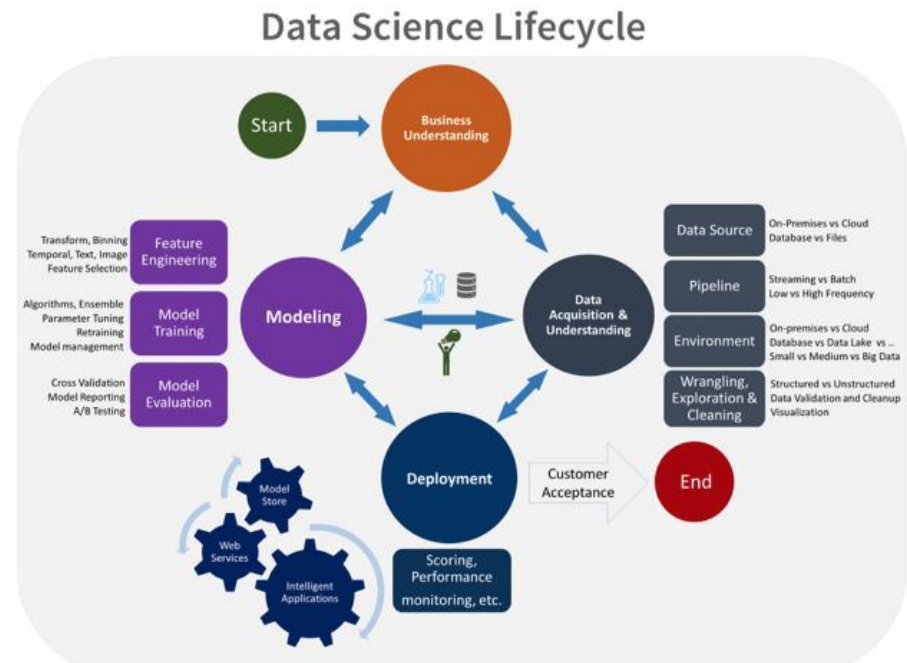
Google's Data Lifecycle

- <https://cloud.google.com/solutions/data-lifecycle-cloud-platform>
- Steps
 - Ingest
 - Store
 - Process and analyze
 - Explore and visualize

Ingest	Store	Process & Analyze	Explore & Visualize
 App Engine	 Cloud Storage	 Cloud Dataflow	 Cloud Datalab
 Compute Engine	 Cloud SQL	 Cloud Dataproc	 Google Data Studio
 Kubernetes Engine	 Cloud Datastore	 BigQuery	 Google Sheets
 Cloud Pub/Sub	 Cloud Bigtable	 Cloud ML	
 Stackdriver Logging	 BigQuery	 Cloud Vision API	
 Cloud Transfer Service	 Cloud Storage for Firebase	 Cloud Speech API	
 Transfer Appliance	 Cloud Firestore	 Translate API	
	 Cloud Spanner	 Cloud Natural Language API	
		 Cloud Dataprep	
		 Cloud Video Intelligence API	

Microsoft Team Data Science Process

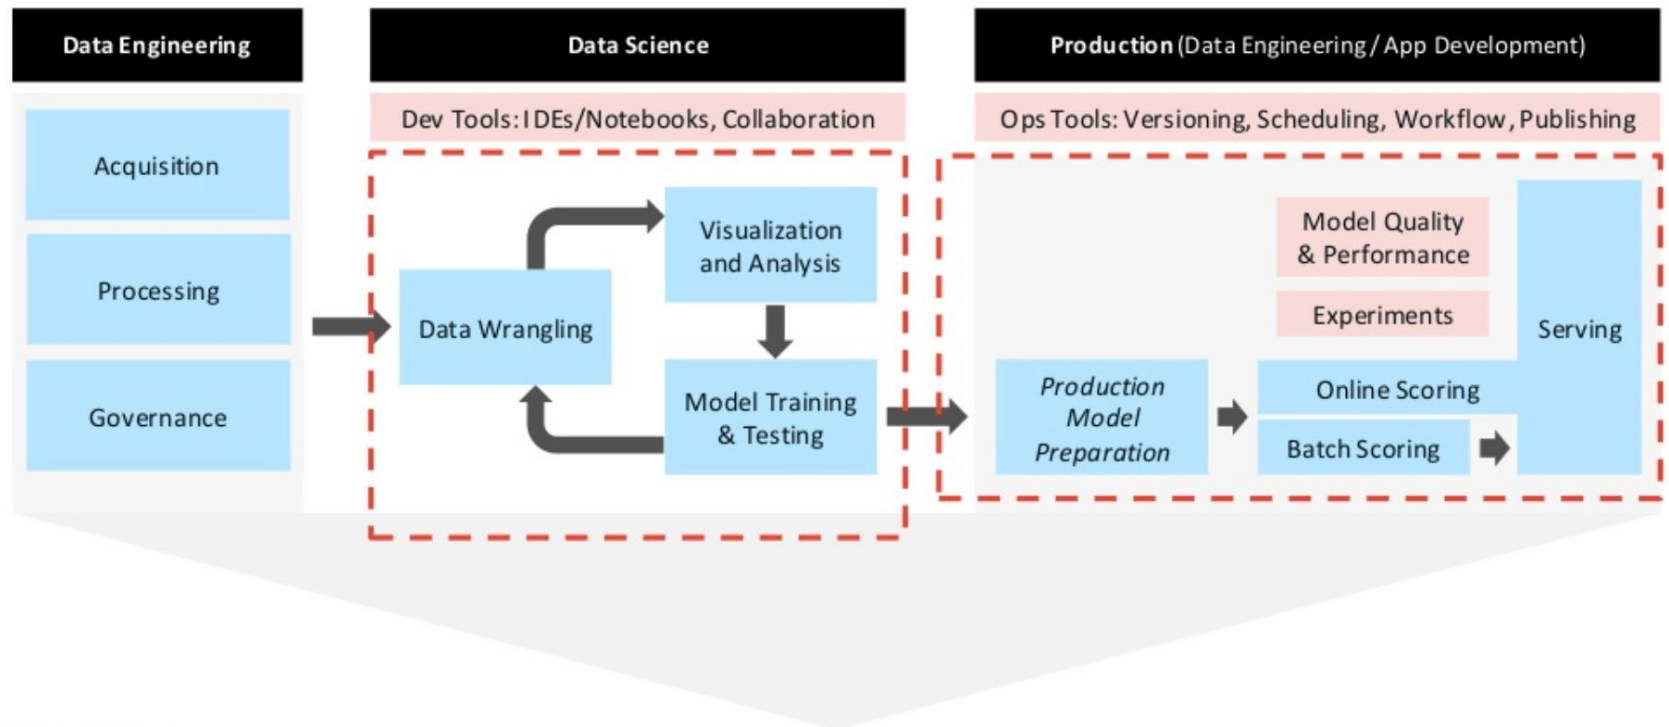
- <https://docs.microsoft.com/en-us/azure/machine-learning/team-data-science-process/overview>
- Steps
 - Business understanding
 - Data acquisition and understanding
 - Modeling
 - Deployment
 - Customer acceptance



Cloudera

Clip slide

End to End Lifecycle of Data Science



Data Analytics Process Approaches

The End

Defining the Problem

Defining the Problem

- Questions to ask
 - What is the business problem you are trying to address?
 - How is it currently done?
 - What performance measure will you use to evaluate your model?
 - What performance criteria do you need to meet to be considered successful?
- Frame it in terms of an analytical solution
 - What kind of machine learning task is it (supervised, unsupervised, classification, regression)?
 - Is it batch or online learning?

Example: The Boston Housing Dataset

- Our target is median housing price
- This is a regression problem
- The first example we did used only a single feature, but this time we are going to use more, so it is a multiple regression problem
- It is known as a univariate regression problem since we are only trying to predict a single value
- It is a fairly small dataset that can fit into memory, so this is a batch learning problem

Possible Performance Measures

- A possible performance measure for regression problems is the root mean square error (RMSE)

$$\sqrt{\frac{1}{m} \sum_{i=1}^m (h(x^{(i)}) - y^{(i)})^2}$$

- Your hypothesis is h
- Your true label is y
- Your number of examples is m
- The i th example is $x^{(i)}$

Mean Absolute Error

- Another possible performance measure for regression is mean absolute error (MAE)
- Might want to use it if there are a lot of outliers in your data
 - Think about it: what happens if you square an outlier

$$\frac{1}{m} \sum_{i=1}^m |h(x^{(i)}) - y^{(i)}|$$

What Is the Difference Between These Two?

- They are distance measures
 - Also called norms
- RMSE is the Euclidean norm, also called the l_2 norm
- MAE is called l_1 norm, the distance travelled along the x and y axis to get to the point of the vector
- If outliers aren't an issue, RMSE is preferred

Defining the Problem

The End

Getting the Data

Getting the Data

- Previously, we loaded the data via a Bunch object in scikit-learn
- Let's now take that Bunch object and make it into a pandas DataFrame
 - This makes it easier to explore the data
- In the scikit-learn Bunch class, which is like a dictionary, here are the following classes:
 - Data
 - Target
 - Feature_names
 - DESCR
- Data, target, and feature_names are ndarrays

Converting Data to a Pandas DataFrame

- Here is the code to convert the data to a NumPy array

```
In [5]: data = pd.DataFrame(boston_dataset.data, columns = boston_dataset.feature_names)
```

```
In [6]: data.head(5)
```

Out[6]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33

Adding the Target to the Pandas DataFrame

- Now let's add the target to this NumPy array

```
In [9]: data['Price'] = boston_dataset.target  
data.head()
```

Out[9]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	Price
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98	24.0
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14	21.6
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03	34.7
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94	33.4
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33	36.2

Looking at the Data Structure

- Try `data.info()` to look at the total number of rows, the attribute's type and the number of non-null values

```
In [16]: data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
CRIM          506 non-null float64
ZN            506 non-null float64
INDUS         506 non-null float64
CHAS          506 non-null float64
NOX           506 non-null float64
RM            506 non-null float64
AGE           506 non-null float64
DIS           506 non-null float64
RAD           506 non-null float64
TAX           506 non-null float64
PTRATIO       506 non-null float64
B             506 non-null float64
LSTAT         506 non-null float64
Price         506 non-null float64
dtypes: float64(14)
memory usage: 55.4 KB
```

Getting the Data

The End

Exploring the Data

Description of the Boston Dataset

- The scikit-learn Bunch class has a DESCR element that tells you about the dataset. (Question: What is the categorical variable?)

:Attribute Information (in order):

- CRIM per capita crime rate by town
- ZN proportion of residential land zoned for lots over 25,000 sq.ft.
- INDUS proportion of non-retail business acres per town
- CHAS Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)
- NOX nitric oxides concentration (parts per 10 million)
- RM average number of rooms per dwelling
- AGE proportion of owner-occupied units built prior to 1940
- DIS weighted distances to five Boston employment centres
- RAD index of accessibility to radial highways
- TAX full-value property-tax rate per \$10,000
- PTRATIO pupil-teacher ratio by town
- B $1000(B_k - 0.63)^2$ where B_k is the proportion of blacks by town
- LSTAT % lower status of the population
- MEDV Median value of owner-occupied homes in \$1000's

Which Feature Is Categorical?

- From the description you should be able to see CHAS is a categorical variable—it is a 1 if it is on the Charles River and 0 otherwise
- Let's see how many homes are on the Charles River:

```
In [15]: data["CHAS"].value_counts()
```

```
Out[15]: 0.0    471  
         1.0     35  
         Name: CHAS, dtype: int64
```

Let's Get a Summary of the Numerical Attributes

- The pandas DataFrame `describe()` shows a summary of the numerical attributes. (Question: What do you notice about the max of the price?)

```
In [12]: data.describe()
```

```
Out[12]:
```

	CRIM	ZN	INDUS	CHAS	NOX
count	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.593761	11.363636	11.136779	0.069170	0.554695
std	8.596783	23.322453	6.860353	0.253994	0.115878
min	0.006320	0.000000	0.460000	0.000000	0.385000
25%	0.082045	0.000000	5.190000	0.000000	0.449000
50%	0.256510	0.000000	9.690000	0.000000	0.538000
75%	3.647423	12.500000	18.100000	0.000000	0.624000
max	88.976200	100.000000	27.740000	1.000000	0.871000

RAD	TAX	PTRATIO	B	LSTAT	Price
3.000000	506.000000	506.000000	506.000000	506.000000	506.000000
9.549407	408.237154	18.455534	356.674032	12.653063	22.532806
3.707259	168.537116	2.164946	91.294864	7.141062	9.197104
1.000000	187.000000	12.600000	0.320000	1.730000	5.000000
4.000000	279.000000	17.400000	375.377500	6.950000	17.025000
5.000000	330.000000	19.050000	391.440000	11.360000	21.200000
4.000000	666.000000	20.200000	396.225000	16.955000	25.000000
4.000000	711.000000	22.000000	396.900000	37.970000	50.000000

Plot Histograms

- A quick visual way to look at the characteristics of your data is to print a histogram of your numerical attributes
- You can see here that Price is capped
- You can also see CHAS is definitely a categorical variable



Create and Hold Out a Test Set

- Use sklearn's `train_test_split` to split data into training and test sets

```
from sklearn.model_selection import train_test_split
train_set, test_set = train_test_split(data, test_size=0.2, random_state=42)
```

```
print(train_set.shape)
print(test_set.shape)
print(type(train_set.shape))
print(train_set[:5])
```

```
(404, 14)
(102, 14)
<class 'tuple'>
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
477	15.02340	0.0	18.10	0.0	0.6140	5.304	97.3	2.1007	24.0	666.0	
15	0.62739	0.0	8.14	0.0	0.5380	5.834	56.5	4.4986	4.0	307.0	
332	0.03466	35.0	6.06	0.0	0.4379	6.031	23.3	6.6407	1.0	304.0	
423	7.05042	0.0	18.10	0.0	0.6140	6.103	85.1	2.0218	24.0	666.0	
19	0.72580	0.0	8.14	0.0	0.5380	5.727	69.5	3.7965	4.0	307.0	

	PTRATIO	B	LSTAT	Price
477	20.2	349.48	24.91	12.0
15	21.0	395.62	8.47	19.9
332	16.9	362.25	7.83	19.4
423	20.2	2.52	23.29	13.4
19	21.0	390.95	11.28	18.2

Exploring the Data

- For regression problems it is important to look at correlation
- This is also called Pearson's r
- Coefficient range from -1 to 1
 - 1 means strong positive correlation
 - -1 means strong negative correlation
 - Close to zero means no linear correlation

Exploring the Data

The End

Visualizing Correlation

Panda's Correlation Method

- Let's look at correlation between our features and target (Price)
- Average number of rooms is highly correlated with price and LSTAT is negatively correlated with price

```
In [28]: corr_matrix = data.corr()
```

```
In [29]: corr_matrix["Price"].sort_values(ascending=False)
```


```
Out[29]: Price      1.000000  
RM        0.695360  
ZN        0.360445  
B         0.333461  
DIS       0.249929  
CHAS      0.175260  
AGE      -0.376955  
RAD      -0.381626  
CRIM     -0.385832  
NOX      -0.427321  
TAX      -0.468536  
INDUS    -0.483725  
PTRATIO  -0.507787  
LSTAT    -0.737663  
Name: Price, dtype: float64
```

Visualize Correlations: sns.heatmap

- A nice way to visualize correlations is with a seaborn heatmap
- Lighter colors have higher correlation
- Darker colors have negative correlation

```
In [41]: import seaborn as sns
```


```
plt.figure(figsize=(10,10))  
ax = sns.heatmap(corr_matrix, annot=True)  
plt.show()
```



Let's Look at Features Highly Correlated With Price

- Average number of rooms: 0.7
- Percentile lower status of the population (LSTAT): -0.74
- Pandas has a `scatter_matrix` or seaborn has a `pairplot`

```
In [47]: sns.pairplot(data[["Price", "LSTAT", "RM"]])  
plt.show()
```



Visualizing Correlation

The End

Preparing the Data

Preparing the Data

- Provide a feature set that does not include the label
- Include a label series to provide the model

```
features=train_set.drop("Price",axis=1)
label=train_set["Price"].copy()
print(type(features))
print(features.head())
print(type(label))
print(label.head())
```

```
<class 'pandas.core.frame.DataFrame'>
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	D
477	15.02340	0.0	18.10	0.0	0.6140	5.304	97.3	2.10
15	0.62739	0.0	8.14	0.0	0.5380	5.834	56.5	4.49
332	0.03466	35.0	6.06	0.0	0.4379	6.031	23.3	6.64
423	7.05042	0.0	18.10	0.0	0.6140	6.103	85.1	2.02
19	0.72580	0.0	8.14	0.0	0.5380	5.727	69.5	3.79

	PTRATIO	B	LSTAT
477	20.2	349.48	24.91
15	21.0	395.62	8.47
332	16.9	362.25	7.83
423	20.2	2.52	23.29
19	21.0	390.95	11.28

```
<class 'pandas.core.series.Series'>
```

477	12.0
15	19.9
332	19.4
423	13.4
19	18.2

```
Name: Price, dtype: float64
```

Deal With Missing Data

- Get rid of attributes with missing data
- Get rid of examples with missing data
- Set missing data to some value
- If you chose option 3 (I'll take the car, Drew) you will need to use an imputer
- A transformer in scikit-learn is an estimator with a fit and then a transform method
- An imputer is an example of a transformer

Handling Categorical Variables

- Recall in our example, CHAS is 1 if it is Charles Riverfront and 0 otherwise
- So the machine learning model doesn't think these categories are ordered, we one-hot encode them so only one value will be equal to 1
- Notice this returns a sparse matrix
- You can turn it back into a NumPy array with the toArray method

```
In [73]: from sklearn.preprocessing import OneHotEncoder
cat_encoder=OneHotEncoder()
CHAS_1hot = cat_encoder.fit_transform(data[["CHAS"]])
CHAS_1hot
```

```
Out[73]: <506x2 sparse matrix of type '<class 'numpy.float64'>'
         with 506 stored elements in Compressed Sparse Row format>
```

```
In [75]: print(CHAS_1hot.toarray())
```

```
[[1. 0.]
 [1. 0.]
 [1. 0.]
 ...
 [1. 0.]
 [1. 0.]
 [1. 0.]]
```


Scaling Numerical Features

- It is important to get all of the input variables in the same scale
- There are two types:
 1. Min-max scaling
 2. Standardization
- Min-max scaling forces everything to between 0 and 1
- Standardization uses the concept of unit variance
- scikit-learn has a `StandardScaler` method to do this

Implementing Transformations on a Dataset

- So we need to one-hot encode our categorical variable and scale our numerical variables
- scikit-learn's ColumnTransformer can help us with this
- It lets you specify which columns to do certain transformations on

```
In [26]: #only numerical features need to be scaled:
num_features = features.drop("CHAS",axis=1)
num_features.columns

Out[26]: Index(['CRIM', 'ZN', 'INDUS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD', 'TAX',
               'PTRATIO', 'B', 'LSTAT'],
              dtype='object')

In [27]: from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.compose import ColumnTransformer
num_attribs = list(num_features)
cat_attribs = ["CHAS"]

prep_pipeline = ColumnTransformer([
    ("std_scaler", StandardScaler(),num_attribs),
    ("one_hot", OneHotEncoder(),cat_attribs)
])

In [28]: data_prepared = prep_pipeline.fit_transform(features)
```

Preparing the Data

The End

Select and Train a Model

Select and Train a Model

- Let's start with a simple linear regressor
- We first fit the model to our prepared data
- We then run the predict method
- Finally, we figure out root mean square error (RMSE) using the known labels on the training set

```
In [53]: from sklearn.linear_model import LinearRegression

lr = LinearRegression()
lr.fit(data_prepared, label)

Out[53]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)

In [36]: some_data = features.iloc[:5]
some_labels = label.iloc[:5]
some_data_prepared = prep_pipeline.transform(some_data)
print("Predictions:", lr.predict(some_data_prepared))

Predictions: [10.96952405 19.41196567 23.06419602 12.1470648 18.3738116 ]

In [37]: print("Labels:", list(some_labels))

Labels: [12.0, 19.9, 19.4, 13.4, 18.2]

In [39]: from sklearn.metrics import mean_squared_error
bh_predictions = lr.predict(data_prepared)
lr_mse = mean_squared_error(label, bh_predictions)
lr_rmse = np.sqrt(lr_mse)
lr_rmse

Out[39]: 4.6520331848801675
```

Comparison With One-Feature Model

- Our RMSE is \$4,600
- Comparing this against the RMSE of our model, just taking average rooms into account you can see we are doing better
- Our RMSE on the initial model is \$6,400

```
In [39]: from sklearn.metrics import mean_squared_error
bh_predictions = lr.predict(data_prepared)
lr_mse = mean_squared_error(label,bh_predictions)
lr_rmse = np.sqrt(lr_mse)
lr_rmse
```

```
Out[39]: 4.6520331848801675
```

```
In [70]: #Let's compare this against our one-feature analysis earlier:
from sklearn.metrics import mean_squared_error
price_room_predictions = price_room.predict(num_Rooms_Train)
pr_mse = mean_squared_error(med_price_Train,price_room_predictions)
pr_rmse = np.sqrt(pr_mse)
pr_rmse
```

```
Out[70]: 6.4034724257657505
```

Let's Try a More Complex Model

- Let's implement a decision tree regressor
- Decision trees are known for discovering non-linear relationships in the data
- Again, this is the process:
 - Fit
 - Train
 - Evaluate
- An RMSE of 0? Do you think we might be overfitting?

```
In [69]: from sklearn.tree import DecisionTreeRegressor
          #print(data_prepared)
          #print(label)
          #print(label.values)

          tree_reg = DecisionTreeRegressor()

          tree_reg.fit(data_prepared, label)
          tree_predictions = tree_reg.predict(data_prepared)
          tree_mse = mean_squared_error(label, tree_predictions)
          tree_rmse = np.sqrt(tree_mse)
          tree_rmse
```

```
Out[69]: 0.0
```

Combating Overfitting: Cross-Validation

- K-fold cross-validation
 - Randomly splits training data into K folds
 - Trains and evaluates K number of times
 - Picks a different fold for evaluation every time while training on the other ones
 - Result is an array containing 10 evaluation scores

Here Is What Cross-Validation of the Decision Tree Shows Us

- Recall our linear regression RMSE on the training set was 4.65

```
In [71]: from sklearn.model_selection import cross_val_score
scores = cross_val_score(tree_reg, data_prepared, label, scoring="neg_mean_squared_error", cv=10)
tree_rmse_scores = np.sqrt(-scores)

In [72]: print("Scores:", tree_rmse_scores)

Scores: [3.8670276  4.3577797  4.80979691 6.73351646 8.09734524 3.60797727
 5.16846205 5.21423532 4.23447163 5.46372126]

In [73]: print("Mean:", tree_rmse_scores.mean())

Mean: 5.155433345174141

In [74]: print("Standard deviation:", tree_rmse_scores.std())

Standard deviation: 1.300090643764176
```

Let's Do Cross-Validation on the Linear Model As Well

```
: from sklearn.model_selection import cross_val_score
lr_scores = cross_val_score(lr, data_prepared, label, scoring="neg_mean_squared_error", cv=10)
lr_rmse_scores = np.sqrt(-lr_scores)
```

```
: print("Scores:", lr_rmse_scores)
```

```
Scores: [3.76298481 4.25110998 5.34719644 6.71464778 4.59265163 5.17395941
 4.43145447 4.5777583  3.6723473  5.77030866]
```

```
: print("Mean:", lr_rmse_scores.mean())
```

```
Mean: 4.829441880454275
```

```
: print("Standard deviation:", lr_rmse_scores.std())
```

```
Standard deviation: 0.8896329730748556
```

Let's Try Another Type of Model

- RandomForestRegressor
 - Works by training many decision trees on random subset of features
 - Averages out their predictions
 - This is an example of ensemble learning, building a model with many other models as input

```
forest_rmse = np.sqrt(forest_mse)
forest_rmse
In[81]: 1.4074152789161882
```

Much better!

```
Scores: [3.00339058 3.37508625 2.65144004 5.38944006 4.63877392 2.91343851
3.78049652 2.56220162 2.59026261 4.47926417]
Mean: 3.538379434860213
Standard Deviation: 0.9446225302809618
```

Select and Train a Model

The End

Tune Your Model

Tune Your Model

- Different models have settings called hyper-parameters that you use to tune the architecture of the models
- Models in scikit-learn have a `get_params` method to show you what these settings are
- Let's look at these for our forest regressor

```
forest_reg.get_params()  
  
{'bootstrap': True,  
 'ccp_alpha': 0.0,  
 'criterion': 'mse',  
 'max_depth': None,  
 'max_features': 'auto',  
 'max_leaf_nodes': None,  
 'max_samples': None,  
 'min_impurity_decrease': 0.0,  
 'min_impurity_split': None,  
 'min_samples_leaf': 1,  
 'min_samples_split': 2,  
 'min_weight_fraction_leaf': 0.0,  
 'n_estimators': 100,  
 'n_jobs': None,  
 'oob_score': False,  
 'random_state': None,  
 'verbose': 0,  
 'warm_start': False}
```

scikit-learn GridSearch Utility

- To explore combinations of hyper-parameters that perform better, you can use scikit-learn's GridSearchCV
- You tell it which hyper-parameters to experiment with
- It will evaluate combinations and values you give it with cross-validation

Import the Constructor and Build a param_grid

- param_grid is a list of dictionary items that define hyper-parameter combinations you want to test
- In this case, we are looking at 12 combinations of n_estimators and max_features and another one with six combinations of those hyper-parameters with bootstrapping off
- The total number of experiments will be 18

```
'warm_start': False}  
  
In [94]: from sklearn.model_selection import GridSearchCV  
         param_grid = [{'n_estimators':[3,10,30], 'max_features':[2,4,6,8]},  
                        {'bootstrap': [False], 'n_estimators':[3,10], 'max_features':[2,3,4]}]
```


Instantiate a grid_search Object

- Inputs include:
 - Your model
 - Your parameter grid
 - The number of folds for cross-validation
 - Your scoring value
- GridSearchCV is like an estimator in the fact that you can fit, predict, score, and transform your data with it

```
In [95]: grid_search = GridSearchCV(forest_reg, param_grid, cv=5, scoring='neg_mean_squared_error',
                                   return_train_score=True)

In [96]: grid_search.fit(data_prepared, label)

Out[96]: GridSearchCV(cv=5, error_score=nan,
                      estimator=RandomForestRegressor(bootstrap=True, ccp_alpha=0.0,
                                                         criterion='mse', max_depth=None,
                                                         max_features='auto',
                                                         max_leaf_nodes=None,
                                                         max_samples=None,
                                                         min_impurity_decrease=0.0,
                                                         min_impurity_split=None,
                                                         min_samples_leaf=1,
                                                         min_samples_split=2,
                                                         min_weight_fraction_leaf=0.0,
                                                         n_estimators=100, n_jobs=None,
                                                         oob_score=False, random_state=None,
                                                         verbose=0, warm_start=False),
                      iid='deprecated', n_jobs=None,
                      param_grid=[{'max_features': [2, 4, 6, 8],
                                    'n_estimators': [3, 10, 30]},
                                   {'bootstrap': [False], 'max_features': [2, 3, 4],
                                    'n_estimators': [3, 10]}],
                      pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
                      scoring='neg_mean_squared_error', verbose=0)
```

Attributes Available for grid_search

- To come up with the best parameter combination use `.best_params_`
- To list the best model use `.best_estimator_`

```
scoring= neg_mean_squared_error , verbose=0),

|: grid_search.best_params_
|: {'max_features': 6, 'n_estimators': 30}

|: grid_search.best_estimator_
|: RandomForestRegressor(bootstrap=True, ccp_alpha=0.0, criterion='mse',
                        max_depth=None, max_features=6, max_leaf_nodes=None,
                        max_samples=None, min_impurity_decrease=0.0,
                        min_impurity_split=None, min_samples_leaf=1,
                        min_samples_split=2, min_weight_fraction_leaf=0.0,
                        n_estimators=30, n_jobs=None, oob_score=False,
                        random_state=None, verbose=0, warm_start=False)
```

Viewing Results

- `.cv_results_` returns a dictionary with keys as column headers and values as columns from across all the hyperparameter combinations
- You can see here the lowest score is 3.63 with 6 `max_features` and 30 `n_estimators`

```
for mean_score, params in zip(cvres["mean_test_score"], cvres["params"]):  
    print(np.sqrt(-mean_score), params)
```

```
4.919806652356437 {'max_features': 2, 'n_estimators': 3}  
4.0001361288564565 {'max_features': 2, 'n_estimators': 10}  
3.9286119149811984 {'max_features': 2, 'n_estimators': 30}  
4.830781872320711 {'max_features': 4, 'n_estimators': 3}  
3.9873853038963256 {'max_features': 4, 'n_estimators': 10}  
3.7762390599094413 {'max_features': 4, 'n_estimators': 30}  
4.705809779306039 {'max_features': 6, 'n_estimators': 3}  
3.8653919909663963 {'max_features': 6, 'n_estimators': 10}  
3.630549635197833 {'max_features': 6, 'n_estimators': 30}  
4.335652220638375 {'max_features': 8, 'n_estimators': 3}  
3.7467842121076 {'max_features': 8, 'n_estimators': 10}  
3.8252166029743706 {'max_features': 8, 'n_estimators': 30}  
4.749045601179783 {'bootstrap': False, 'max_features': 2, 'n_estimators': 3}  
3.9302707749209254 {'bootstrap': False, 'max_features': 2, 'n_estimators': 10}  
4.713771588928827 {'bootstrap': False, 'max_features': 3, 'n_estimators': 3}  
3.7404543607374685 {'bootstrap': False, 'max_features': 3, 'n_estimators': 10}  
4.725776175352956 {'bootstrap': False, 'max_features': 4, 'n_estimators': 3}  
3.8311000092485443 {'bootstrap': False, 'max_features': 4, 'n_estimators': 10}
```

Inspect Your Results

- Which features are important?
- This will show it for you:
 `.best_estimator_.feature_importances_`
- Let's look at a sorted list of the most important features
 - We need to provide a list of numerical attributes and one-hot encoded categories in order to do this
- Not surprisingly, notice the most important features were those with high correlation in our initial analysis
- We could use this information to implement a simpler model

```
|: feature_importances = grid_search.best_estimator_.feature_importances_  
  
|: cat_encoder = prep_pipeline.named_transformers_["one_hot"]  
   cat_one_hot_attribs = list(cat_encoder.categories_[0])  
   cat_one_hot_attribs  
  
|: [0.0, 1.0]  
  
|: sorted(zip(feature_importances, (num_attribs+cat_one_hot_attribs)))  
|: [(0.003059841389540117, 0.0),  
   (0.006072988585918179, 'ZN'),  
   (0.007010766272373032, 'RAD'),  
   (0.008358714641886356, 1.0),  
   (0.01839982084135831, 'B'),  
   (0.018466831659624712, 'TAX'),  
   (0.020872801565644968, 'AGE'),  
   (0.030008410456314357, 'PTRATIO'),  
   (0.034131419179445614, 'CRIM'),  
   (0.05560177643879219, 'DIS'),  
   (0.05987860450107494, 'INDUS'),  
   (0.06084123851005931, 'NOX'),  
   (0.33430816725420187, 'RM'),  
   (0.3429886187037661, 'LSTAT')]
```

Simpler = better

Evaluate Your Best Model on Your Test Set

- You need to perform the same data prep on the test set to feed it to your model
- To prepare your data, remember to do just a transform, not a fit_transform
- From there, it is the same evaluation steps you did on your training data

```
] final_model = grid_search.best_estimator_  
  
X_test = test_set.drop("Price", axis=1)  
y_test = test_set["Price"].copy()  
  
] X_test_prepared = prep_pipeline.transform(X_test)  
final_predictions = final_model.predict(X_test_prepared)  
final_mse = mean_squared_error(y_test, final_predictions)  
final_rmse = np.sqrt(final_mse)  
final_rmse  
] 3.0343068178619887
```

Save Your Model for Future Use

- You can use Python's PICKLE object serializer to save your model
- You can also use joblib
 - It has a dump method to serialize and save your model
 - It has a load method to read it back in and make it usable

```
from sklearn.externals import joblib
joblib.dump(final_model, "my_model.pkl")
```

```
/Users/rjdaskevich/anaconda3/lib/python3.6
: sklearn.externals.joblib is deprecated in
directly from joblib, which can be install
ickled models, you may need to re-serializ
warnings.warn(msg, category=FutureWarnin
```

```
['my_model.pkl']
```

```
!ls my_model.*
```

```
my_model.pkl
```

Tune Your Model

The End

Present Your Solution

Present Your Solution

- What have you learned?
- What worked?
- What didn't work?
- What assumptions were made?
- What limitations did you have?
- What are your final recommendations?
- Don't forget to provide visualizations and clear results

Launch Production System

- Hook your system up to production input data
- Check prediction performance at regular intervals
- Look at input data quality
- Make sure you schedule time to retrain your models regularly

Present Your Solution

The End