

# COMP 4432 Machine Learning

---

## Lesson 4: Model Training

# Agenda

---

- Training Models
  - Theory and Implementation
  - Requires an understanding an algorithm's capabilities and hyperparameters
  - When somethings breaks, know and recognize what is at fault

# The Process

---

Objective

Gathered &  
Prepared  
Data

Identified  
Potential  
Algorithms

Model  
Training and  
Selection

Production

# The Process

---

Objective

Gathered &  
Prepared  
Data

Identified  
Potential  
Algorithms

Model  
Training and  
Selection

Production

# Training Models

---

- Easy to execute
- Easy to mishandle
  - Generalizability
- More to it than just “*.fit()*”

# Training Models

---

- Easy to execute
- Easy to mishandle
  - Generalizability
- More to it than just “*.fit()*”
- New tools available to “simplify” training
  - AutoML (Automated Machine Learning)
    - FLAML, AutoGluon, AutoPilot
  - Hyperparameter Tuning Optimization
    - (HyperOpt, OpTuna)

# Training Models

---

- When “*.fit()*” is called, what is happening?
- Finding optimum model parameters by minimizing a cost function

$$J(\theta) = \frac{1}{2n} \sum_{i=1}^n (h_{\theta}(x_i) - y_i)^2$$
$$h_{\theta}(x^{(i)}) = \theta_0 x_0^{(i)} + \theta_1 x_1^{(i)} + \dots + \theta_f x_f^{(i)} \quad h_{\theta}(x^{(i)}) = \theta^T x^{(i)}$$

$$J(\theta) = -\frac{1}{n} \sum_i y_i \log(p(x_i)) + (1 - y_i) \log(1 - p(x_i))$$
$$p = \frac{1}{1 + e^{\vec{\beta} \cdot \vec{x}}}$$

$$J(k, t_k) = \frac{m_{\text{left}}}{m} \text{MSE}_{\text{left}} + \frac{m_{\text{right}}}{m} \text{MSE}_{\text{right}} \quad \text{where} \quad \begin{cases} \text{MSE}_{\text{node}} = \sum_{i \in \text{node}} (\hat{y}_{\text{node}} - y^{(i)})^2 \\ \hat{y}_{\text{node}} = \frac{1}{m_{\text{node}}} \sum_{i \in \text{node}} y^{(i)} \end{cases}$$

# Training Models

---

- When “*.fit()*” is called, what is happening?
- Finding optimum model parameters by minimizing a cost function
  - Theta (Beta) vector in linear (logistic) regression
  - Features and splits in a decision tree



# Training Models

---

- When “*.fit()*” is called, what is happening?
- Finding optimum model parameters by minimizing a cost function
  - Theta (Beta) vector in linear (logistic) regression
  - Features and splits in a decision tree
- Dependent on user-defined hyperparameters
  - [Linear Regression](#)
  - [Random Forest Regression](#)

# Linear Regression (Closed Form)

---

$$J(\theta) = \frac{1}{n} \sum_{i=1}^n \left( h_{\theta}(x^{(i)}) - y^{(i)} \right)^2$$

$$h_{\theta}(x^{(i)}) = \theta_0 x_0^{(i)} + \theta_1 x_1^{(i)} + \dots + \theta_f x_f^{(i)} \quad h_{\theta}(x^{(i)}) = \theta^T x^{(i)}$$

# Linear Regression (Closed Form)

---

$$J(\theta) = \frac{1}{n} \sum_{i=1}^n \left( h_{\theta}(x^{(i)}) - y^{(i)} \right)^2$$

$$h_{\theta}(x^{(i)}) = \theta_0 x_0^{(i)} + \theta_1 x_1^{(i)} + \dots + \theta_f x_f^{(i)} \quad h_{\theta}(x^{(i)}) = \theta^T x^{(i)}$$

$$\theta = \begin{pmatrix} \theta_0 \\ \theta_1 \\ \vdots \\ \theta_f \end{pmatrix} \quad x^{(i)} = \begin{pmatrix} x_0^{(i)} \\ x_1^{(i)} \\ \vdots \\ x_f^{(i)} \end{pmatrix}$$

$$J(\theta) = \frac{1}{n} \sum_{i=1}^n \left( \theta^T x^{(i)} - y^{(i)} \right)^2$$

# Linear Regression (Matrix Form)

---

$$J(\theta) = \frac{1}{n} \sum_{i=1}^n \left( \theta^T x^{(i)} - y^{(i)} \right)^2$$

$$X = \begin{pmatrix} x^{(1)T} \\ x^{(2)T} \\ \vdots \\ x^{(n)T} \end{pmatrix} = \begin{pmatrix} x_0^{(1)} & x_1^{(1)} & \dots & x_f^{(1)} \\ x_0^{(2)} & x_1^{(2)} & \dots & x_f^{(2)} \\ \vdots & \vdots & \ddots & \vdots \\ x_0^{(n)} & x_1^{(n)} & \dots & x_f^{(n)} \end{pmatrix} \quad \theta = \begin{pmatrix} \theta_0 \\ \theta_1 \\ \vdots \\ \theta_f \end{pmatrix}$$

# Linear Regression (Matrix Form)

---

$$J(\theta) = \frac{1}{n} \sum_{i=1}^n \left( \theta^T x^{(i)} - y^{(i)} \right)^2$$

$$X = \begin{pmatrix} x^{(1)T} \\ x^{(2)T} \\ \vdots \\ x^{(n)T} \end{pmatrix} = \begin{pmatrix} x_0^{(1)} & x_1^{(1)} & \dots & x_f^{(1)} \\ x_0^{(2)} & x_1^{(2)} & \dots & x_f^{(2)} \\ \vdots & \vdots & \ddots & \vdots \\ x_0^{(n)} & x_1^{(n)} & \dots & x_f^{(n)} \end{pmatrix} \quad \theta = \begin{pmatrix} \theta_0 \\ \theta_1 \\ \vdots \\ \theta_f \end{pmatrix}$$

$$X\theta = \begin{pmatrix} \theta_0 x_0^{(1)} + \theta_1 x_1^{(1)} + \dots + \theta_f x_f^{(1)} \\ \theta_0 x_0^{(2)} + \theta_1 x_1^{(2)} + \dots + \theta_f x_f^{(2)} \\ \vdots \\ \theta_0 x_0^{(n)} + \theta_1 x_1^{(n)} + \dots + \theta_f x_f^{(n)} \end{pmatrix}$$

# Linear Regression (Matrix Form)

---

$$J(\theta) = \frac{1}{n} \sum_{i=1}^n \left( \theta^T x^{(i)} - y^{(i)} \right)^2$$

$$X = \begin{pmatrix} x^{(1)T} \\ x^{(2)T} \\ \vdots \\ x^{(n)T} \end{pmatrix} = \begin{pmatrix} x_0^{(1)} & x_1^{(1)} & \dots & x_f^{(1)} \\ x_0^{(2)} & x_1^{(2)} & \dots & x_f^{(2)} \\ \vdots & \vdots & \ddots & \vdots \\ x_0^{(n)} & x_1^{(n)} & \dots & x_f^{(n)} \end{pmatrix} \quad \theta = \begin{pmatrix} \theta_0 \\ \theta_1 \\ \vdots \\ \theta_f \end{pmatrix}$$

$$X\theta = \begin{pmatrix} \theta_0 x_0^{(1)} + \theta_1 x_1^{(1)} + \dots + \theta_f x_f^{(1)} \\ \theta_0 x_0^{(2)} + \theta_1 x_1^{(2)} + \dots + \theta_f x_f^{(2)} \\ \vdots \\ \theta_0 x_0^{(n)} + \theta_1 x_1^{(n)} + \dots + \theta_f x_f^{(n)} \end{pmatrix}$$

$$J(\theta) = \frac{1}{n} (X\theta - y)^T (X\theta - y)$$

# Linear Regression (Matrix Form)

---

$$J(\theta) = \frac{1}{n} \sum_{i=1}^n \left( \theta^T x^{(i)} - y^{(i)} \right)^2$$

$$J(\theta) = \frac{1}{n} (X\theta - y)^T (X\theta - y)$$

$$J(\theta) = \theta^T X^T X \theta - 2(X\theta)^T y + y^T y$$

# Linear Regression (Matrix Form)

---

$$J(\theta) = \frac{1}{n} \sum_{i=1}^n \left( \theta^T x^{(i)} - y^{(i)} \right)^2$$

$$J(\theta) = \frac{1}{n} (X\theta - y)^T (X\theta - y)$$

$$J(\theta) = \theta^T X^T X \theta - 2(X\theta)^T y + y^T y$$

$$\frac{\partial J}{\partial \theta} = 2X^T X \theta - 2X^T y = 0$$

$$X^T X \theta = X^T y$$

$$\theta = (X^T X)^{-1} X^T y$$



# Linear Regression

---

- Closed form solution is computationally demanding as the number of rows increases
  - Specifically, the inverse

$$\theta = (X^T X)^{-1} X^T y$$

- Gradient Descent offers a solution to large n

$$\theta_j \leftarrow \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

$$\frac{\partial J(\theta)}{\partial \theta_j} = \frac{2}{n} \sum_{i=1}^n \left( \theta^T x^{(i)} - y^{(i)} \right) x_j^{(i)}$$

# Gradient Descent

---

- Start with random values in  $\theta$
- Identify the direction that offers the largest reduction in error with respect to the parameters we are trying to find  $\theta$
- Update the values of  $\theta$  and try again

$$J(\theta) = \frac{1}{n} \sum_{i=1}^n \left( \theta^T x^{(i)} - y^{(i)} \right)^2$$

$$\frac{\partial J(\theta)}{\partial \theta_j} = \frac{2}{n} \sum_{i=1}^n \left( \theta^T x^{(i)} - y^{(i)} \right) x_j^{(i)}$$

$$\theta_j \leftarrow \theta_j - \alpha \frac{\partial J(\theta)}{\partial \theta_j}$$

# Gradient Descent

---

$$J(\theta) = \frac{1}{n} \sum_{i=1}^n \left( \theta^T x^{(i)} - y^{(i)} \right)^2$$

$$\begin{aligned} \theta^T x^{(i)} &= \theta_0 x_0^{(i)} + \theta_1 x_1^{(i)} + \dots + \theta_f x_f^{(i)} \\ \theta^T x^{(i)} - y^{(i)} &= \theta_0 x_0^{(i)} + \theta_1 x_1^{(i)} + \dots + \theta_f x_f^{(i)} - y^{(i)} \end{aligned}$$

# Gradient Descent

---

$$J(\theta) = \frac{1}{n} \sum_{i=1}^n \left( \theta^T x^{(i)} - y^{(i)} \right)^2$$

$$\begin{aligned} \theta^T x^{(i)} &= \theta_0 x_0^{(i)} + \theta_1 x_1^{(i)} + \dots + \theta_f x_f^{(i)} \\ \theta^T x^{(i)} - y^{(i)} &= \theta_0 x_0^{(i)} + \theta_1 x_1^{(i)} + \dots + \theta_f x_f^{(i)} - y^{(i)} \end{aligned}$$

$$\frac{\partial \left( \theta^T x^{(i)} - y^{(i)} \right)^2}{\partial \theta_j} = \frac{\partial \left( \theta_0 x_0^{(i)} + \theta_1 x_1^{(i)} + \dots + \theta_f x_f^{(i)} - y^{(i)} \right)^2}{\partial \theta_j}$$

# Gradient Descent

---

$$J(\theta) = \frac{1}{n} \sum_{i=1}^n \left( \theta^T x^{(i)} - y^{(i)} \right)^2$$

$$\begin{aligned} \theta^T x^{(i)} &= \theta_0 x_0^{(i)} + \theta_1 x_1^{(i)} + \dots + \theta_f x_f^{(i)} \\ \theta^T x^{(i)} - y^{(i)} &= \theta_0 x_0^{(i)} + \theta_1 x_1^{(i)} + \dots + \theta_f x_f^{(i)} - y^{(i)} \end{aligned}$$

$$\begin{aligned} \frac{\partial \left( \theta^T x^{(i)} - y^{(i)} \right)^2}{\partial \theta_j} &= \frac{\partial \left( \theta_0 x_0^{(i)} + \theta_1 x_1^{(i)} + \dots + \theta_f x_f^{(i)} - y^{(i)} \right)^2}{\partial \theta_j} \\ &= 2 \left( \theta_0 x_0^{(i)} + \theta_1 x_1^{(i)} + \dots + \theta_f x_f^{(i)} - y^{(i)} \right) x_j^{(i)} \end{aligned}$$

$$\frac{\partial J(\theta)}{\partial \theta_j} = \frac{2}{n} \sum_{i=1}^n \left( \theta^T x^{(i)} - y^{(i)} \right) x_j^{(i)}$$

# Gradient Descent

---

$$\frac{\partial J(\theta)}{\partial \theta_j} = \frac{2}{n} \sum_{i=1}^n \left( \theta^T x^{(i)} - y^{(i)} \right) x_j^{(i)}$$

$$\begin{aligned} \frac{\partial J(\theta)}{\partial \theta_{j=j'}} = & (\theta_0 x_0^{(1)} + \theta_1 x_1^{(1)} + \dots + \theta_f x_f^{(1)} - y^{(1)}) x_{j'}^{(1)} \\ & + \\ & (\theta_0 x_0^{(2)} + \theta_1 x_1^{(2)} + \dots + \theta_f x_f^{(2)} - y^{(2)}) x_{j'}^{(2)} \\ & + \\ & \dots \\ & + \\ & (\theta_0 x_0^{(n)} + \theta_1 x_1^{(n)} + \dots + \theta_f x_f^{(n)} - y^{(n)}) x_{j'}^{(n)} \end{aligned}$$

# Gradient Descent

---

$$\frac{\partial J(\theta)}{\partial \theta_j} = \frac{1}{n} \sum_{i=1}^n \left( h_{\theta}(x^{(i)}) - y^{(i)} \right) x_j^{(i)}$$

$$\begin{aligned} \frac{\partial J(\theta)}{\partial \theta_{j=j'}} = & (\theta_0 x_0^{(1)} + \theta_1 x_1^{(1)} + \dots + \theta_f x_f^{(1)} - y^{(1)}) x_{j'}^{(1)} \\ & + \\ & (\theta_0 x_0^{(2)} + \theta_1 x_1^{(2)} + \dots + \theta_f x_f^{(2)} - y^{(2)}) x_{j'}^{(2)} \\ & + \\ & \dots \\ & + \\ & (\theta_0 x_0^{(n)} + \theta_1 x_1^{(n)} + \dots + \theta_f x_f^{(n)} - y^{(n)}) x_{j'}^{(n)} \end{aligned}$$

$$X^T = \begin{pmatrix} x_0^{(1)} & x_0^{(2)} & \dots & x_0^{(n)} \\ x_1^{(1)} & x_1^{(2)} & \dots & x_1^{(n)} \\ \vdots & \vdots & \ddots & \vdots \\ x_f^{(1)} & x_f^{(2)} & \dots & x_f^{(n)} \end{pmatrix} \quad X\theta - y = \begin{pmatrix} \theta_0 x_0^{(1)} + \theta_1 x_1^{(1)} + \dots + \theta_f x_f^{(1)} - y^{(1)} \\ \theta_0 x_0^{(2)} + \theta_1 x_1^{(2)} + \dots + \theta_f x_f^{(2)} - y^{(2)} \\ \vdots \\ \theta_0 x_0^{(n)} + \theta_1 x_1^{(n)} + \dots + \theta_f x_f^{(n)} - y^{(n)} \end{pmatrix}$$

# Gradient Descent

---

$$\frac{\partial J(\theta)}{\partial \theta_j} = \frac{1}{n} \sum_{i=1}^n \left( h_{\theta}(x^{(i)}) - y^{(i)} \right) x_j^{(i)}$$

$$\begin{aligned} \frac{\partial J(\theta)}{\partial \theta_{j=j'}} = & (\theta_0 x_0^{(1)} + \theta_1 x_1^{(1)} + \dots + \theta_f x_f^{(1)} - y^{(1)}) x_{j'}^{(1)} \\ & + \\ & (\theta_0 x_0^{(2)} + \theta_1 x_1^{(2)} + \dots + \theta_f x_f^{(2)} - y^{(2)}) x_{j'}^{(2)} \\ & + \\ & \dots \\ & + \\ & (\theta_0 x_0^{(n)} + \theta_1 x_1^{(n)} + \dots + \theta_f x_f^{(n)} - y^{(n)}) x_{j'}^{(n)} \end{aligned}$$

$$X^T = \begin{pmatrix} x_0^{(1)} & x_0^{(2)} & \dots & x_0^{(n)} \\ x_1^{(1)} & x_1^{(2)} & \dots & x_1^{(n)} \\ \vdots & \vdots & \ddots & \vdots \\ x_f^{(1)} & x_f^{(2)} & \dots & x_f^{(n)} \end{pmatrix} \quad X\theta - y = \begin{pmatrix} \theta_0 x_0^{(1)} + \theta_1 x_1^{(1)} + \dots + \theta_f x_f^{(1)} - y^{(1)} \\ \theta_0 x_0^{(2)} + \theta_1 x_1^{(2)} + \dots + \theta_f x_f^{(2)} - y^{(2)} \\ \vdots \\ \theta_0 x_0^{(n)} + \theta_1 x_1^{(n)} + \dots + \theta_f x_f^{(n)} - y^{(n)} \end{pmatrix}$$

$$\nabla J(\theta) = X^T (X\theta - y)$$

$$\theta \leftarrow \theta - \eta \nabla J(\theta)$$



# Regularization

---

- Regularization introduces a penalty for large values of theta
- Capable of eliminating features to address overfitting

$$\theta_j \leftarrow \theta_j - \alpha \frac{\partial J(\theta)}{\partial \theta_j}$$

$$\frac{\partial J(\theta)}{\partial \theta_j} = \frac{2}{n} \sum_{i=1}^n \left( \theta^T x^{(i)} - y^{(i)} \right) x_j^{(i)}$$

# Regularization

---

- Regularization introduces a penalty for large values of theta
- Capable of eliminating features to address overfitting

$$\theta_j \leftarrow \theta_j - \alpha \frac{\partial J(\theta)}{\partial \theta_j} \qquad \frac{\partial J(\theta)}{\partial \theta_j} = \frac{2}{n} \sum_{i=1}^n \left( \theta^T x^{(i)} - y^{(i)} \right) x_j^{(i)}$$

$$J(\theta) = \frac{1}{n} \left[ \sum_{i=1}^n \left( \theta^T x^{(i)} - y^{(i)} \right)^2 + \lambda \sum_{j=1}^f \theta_j^2 \right]$$

$$\theta_j \leftarrow \theta_j \left( 1 - \alpha \frac{\lambda}{n} \right) - \frac{\alpha}{n} \sum_{i=1}^n \left( \theta^T x^{(i)} - y^{(i)} \right) x_j^{(i)}$$

# Regularization

---

- Worked Example

# Polynomial Regression


---

$$f(x) = \theta_0 + \theta_1 x + \theta_2 x^2 + \dots + \theta_j x^j$$

- Typically, we don't know  $j$  beforehand
- `sklearn.linear_model.LinearRegression` still works
- Worked Example

# Learning Curves

- Example from FLAML



# Early Stopping

---

- Definition in XGBoost (`early_stopping_rounds`):  
Validation metric needs to improve at least once in every **`early_stopping_rounds`** rounds to continue training.

# Early Stopping

---

- Definition in XGBoost (`early_stopping_rounds`):  
Validation metric needs to improve at least once in every **early\_stopping\_rounds** rounds to continue training.

